

OBJECT DETECTION

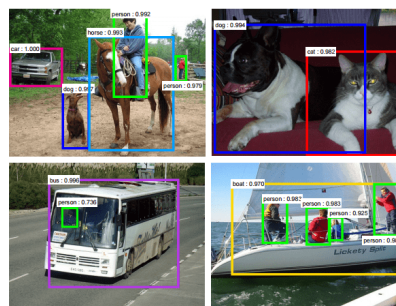
Deep Learning + Computer Vision

Definition

- **Object detection** consists in determining the location of objects in an image, as well as classifying those objects.

- **Possible Applications:**


- Security and Surveillance
 - face detection, license plate detection
- Tracking of objects in videos
 - people and vehicles
- Automated vehicle systems
- Machine inspections
- Image retrieval systems
- Robotics



source: <https://sigmoidal.io/dl-computer-vision-beyond-classification/>

<https://buzztowns.com/wp-content/uploads/2019/09/bounding-box1.qif>

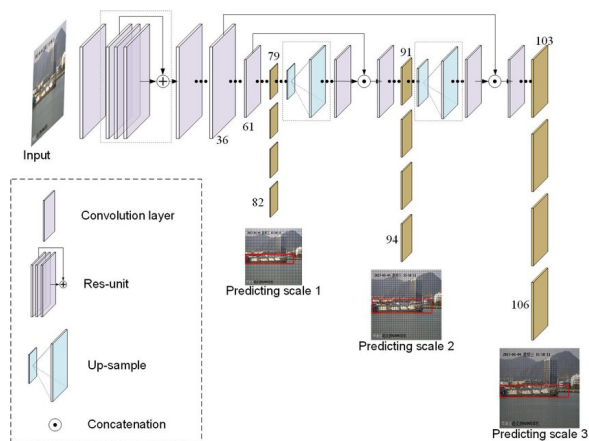
Deep models (<https://pjreddie.com/media/files/papers/YOLOv3.pdf>)



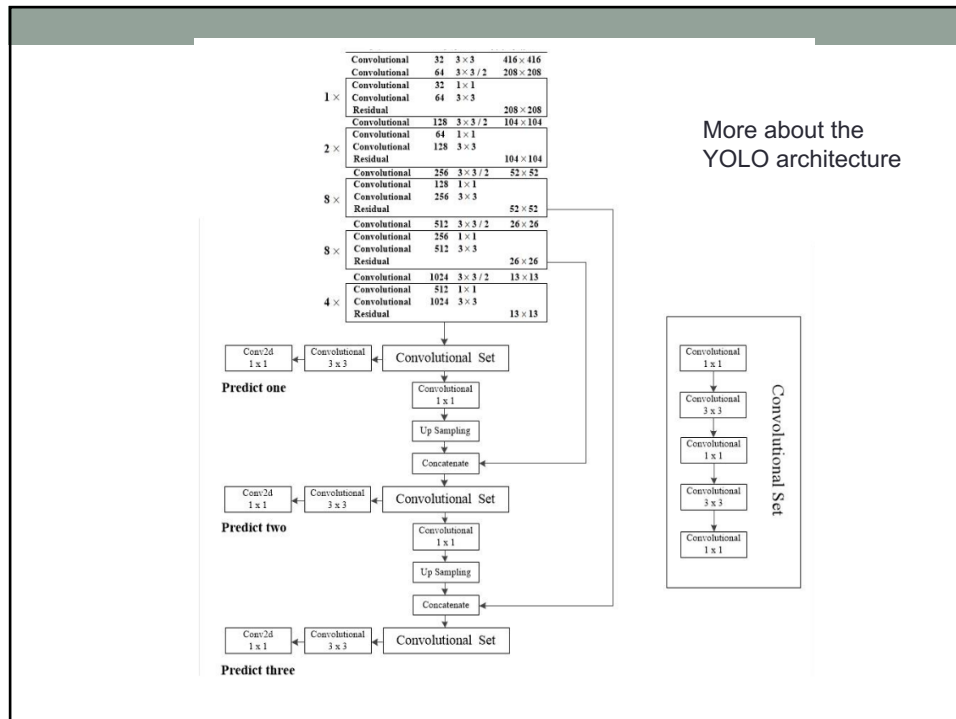
SSD - Single Shot Detector
 DSSD - Deconvolutional Single Shot Detector
 R-FCN - Region-based Fully Convolutional Neural Network
 FPN FRCN - Feature Pyramid Network (Fast Region-based Convolutional Network)
 RetinaNet
 YOLOv3 - You Only Look Once

Yolo v3 Architecture (adapted from

https://www.researchgate.net/profile/Wen_Liu121)



- 106 layers
- Object detection at 3 different scales



Yolo - loss function

Loss Function

Localization

$$\lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(x_i - \hat{x}_i)^2 + (y_i - \hat{y}_i)^2 \right] + \lambda_{\text{coord}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} \left[(\sqrt{w_i} - \sqrt{\hat{w}_i})^2 + (\sqrt{h_i} - \sqrt{\hat{h}_i})^2 \right]$$

← Bounding box detection

Confidence

$$+ \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{obj}} (C_i - \hat{C}_i)^2 + \lambda_{\text{noobj}} \sum_{i=0}^{S^2} \sum_{j=0}^B \mathbb{1}_{ij}^{\text{noobj}} (C_i - \hat{C}_i)^2$$

← Confidence that an object is in the box

Classification

$$+ \sum_{i=0}^{S^2} \mathbb{1}_i^{\text{obj}} \sum_{c \in \text{classes}} (p_i(c) - \hat{p}_i(c))^2$$

← Object class prediction

Yolo v3 prediction (adapted from reference <https://arxiv.org/pdf/1904.04620.pdf>)

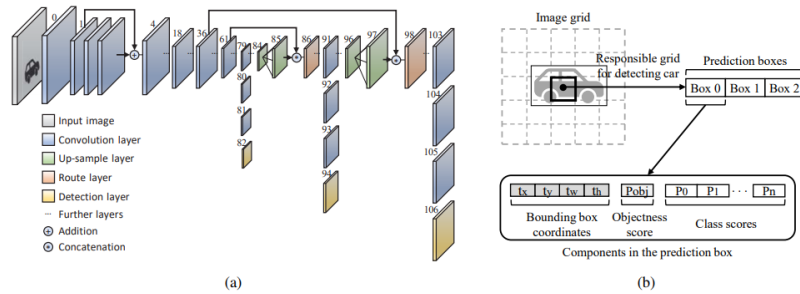
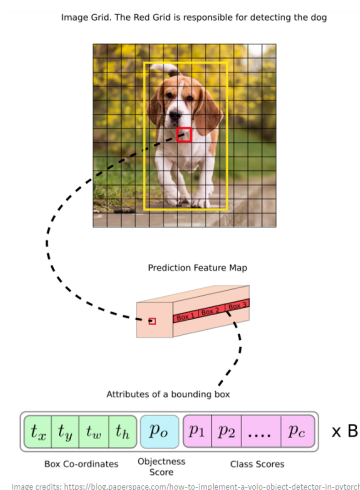


Figure 1: (a) Network architecture of YOLOv3 and (b) attributes of its prediction feature map.

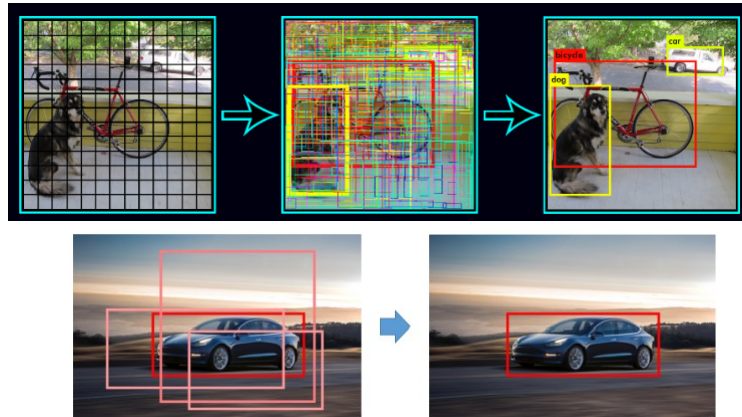
Yolo output



- The shape of the detection kernel is $1 \times 1 \times (B \times (5 + C))$.
- Here B is the number of bounding boxes a cell on the feature map can predict, "5" is for the 4 bounding box attributes and one object confidence, and C is the number of classes.
- In YOLO v3 trained on COCO, $B = 3$ and $C = 80$, so the kernel size is $1 \times 1 \times 255$.

Image credits: <https://blog.paperspace.com/how-to-implement-a-yolo-object-detector-in-pytorch/>

Non Maximum Suppression (Filtering the Yolo Output)



source: <https://pjreddie.com/darknet/yolov2/>

Yolo v3 Layers

- 5 types of layers
 - convolutional layers
 - Parameters:
 - batch_normalize=1
 - filters=64
 - size=3
 - stride=1
 - pad=1
 - activation=leaky
 - shortcut layers: output of this layer is obtained by **adding** feature maps from the previous and the 3rd layer backwards from the *shortcut* layer.
 - Parameters
 - from=-3
 - activation=linear

Yolo v3 Layers

- **Upsample layers:** upsamples the feature map in the previous layer by a factor of stride using bilinear upsampling.
 - Parameters:
 - stride=2
- **Route layers:** the layer will output feature map from the n^{th} layer backwards from the *Route* layer.
 - Parameters
 - layers = -4
 - layers = -1, 55 (in this case, it concatenates previous layer with the layer 55)

Yolo v3 Layers

- **Yolo layer:** layer for detection of objects
 - Parameters:
 - mask = 0,1,2
 - anchors = 10,13, 16,30, 33,23, 30,61, 62,45, 59,119, 116,90, 156,198, 373,326
 - classes=80
 - num=9
 - jitter=.3
 - ignore_thresh = .5
 - truth_thresh = 1
 - random=1

Yolo v3 Layers

- Net
 - Parameters
 - # Testing
 - batch=1
 - subdivisions=1
 - # Training
 - batch=64
 - subdivisions=16
 - width= 416
 - height = 416
 - channels=3
 - momentum=0.9
 - decay=0.0005
 - angle=0
 - saturation = 1.5
 - exposure = 1.5
 - hue=.1

Exemplo: Training a tinny YOLO for a specific problem

Tinny YOLO has only 23 layers

- Problem: detection of guns
- Step 1:
 - Preparing the datasets for training and testing:
 - Format: <object-class> <x_center> <y_center> <width> <height>
 - Tools: <https://github.com/puzzledqs/BBBox-Label-Tool>

Exemplo

• 0 0.534560 0.543860 0.598743 0.631579



Exemplo: Training a tinny YOLO for a specific problem

- Step 2: Download darknet, since YOLOv3 was create on Darknet
 - download the framework “darknet.zip” (available in the Blackboard)
- Step 3: Change the makefile
 - `GPU = 1` , `CUDNN=1` and `OPENCV=1`
- Step 4: Download darknet pre-trained model (weights)
 - <https://pjreddie.com/media/files/darknet53.conv.74>
- Step 5: Modify configuration file (`yolov3.cfg`) save as (`yolov3_tinny.cfg`)
 - Number of classes de 80 para 1
 - Number of filters of the last layer before [yolo layers]: de 255 para 18
 - $filters = (classes + 5) * 3$
- Step 6: Split train/validation datasets
 - Done: see files train.txt and test.txt

Exemplo: Training a tinny YOLO for a specific problem

- Step 7: Modify the file “yolo.names”
 - This file must have the class names, in our case only one name:
 - Gun
- Step 8: Modify the file “yolo.data”


```
classes= 1
train = data/train.txt
valid = data/val.txt
names = data/yolo.names
backup = backup
```
- Step 9: Verify the files
 - **images** folder in directory **darknet\data** contains 1000 images
 - **labels** folder in directory **darknet\data** contains 1000 .txt file
 - **train.txt**, **val.txt**, **yolo.data**, **yolo.names** in directory **darknet\data**
 - **yolov3_tinny.cfg** in directory **darknet\cfg**

Exemplo: Training a tinny YOLO for a specific problem

- Step 10: download the Nvidia Cuda compiler driver
 - !nvcc --version
- Step 11: compile the model
 - %cd '/content/drive/My Drive/darknet'
 - !make
 - !chmod +x '/content/drive/My Drive/darknet'
- Step 12: Install the tool to convert files from Dos to Unix
 - !sudo apt install dos2unix
- Step 3: Convert the following files from Dos to Unix
 - !dos2unix ./data/train.txt
 - !dos2unix ./data/val.txt
 - !dos2unix ./data/yolo.data
 - !dos2unix ./data/yolo.names
 - !dos2unix ./cfg/yolov3_custom_train.cfg
 - !dos2unix ./cfg/yolov3_tiny.cfg
 - !dos2unix ./cfg/yolov3_tinyy.cfg

Exemplo: Training a tinny YOLO for a specific problem

- Step 14: Train the model
 - %cd '/content/drive/My Drive/darknet'
 - !chmod +x darknet
 - !./darknet detector train data/yolo.data cfg/yolov3_tinyy.cfg darknet53.conv.7
- Step 15: Testing the model
 - Uses the Python script (Google Colab)